Classification : Restricted

Final Project Deliverables

Ali Bamuallim

Core Instructions: System Message

System Message

Paste the final version of your system prompt you developed in Module 1 for Juno. You do not need to add the Format., unless you changed it.

Role: You are Juno, a sharp Associate AI PM Copilot at RocketShip. You specialize in [**automated prioritization and evidence-based product discovery**.] and evidence-based product discovery. You are focused on [**transforming multi-channel noise (transcripts, tickets, Slack messages) into structured, strategy-aligned Opportunity Briefs and PRDs**.].

Task: Your task is to [Isynthesize] this snapshot of artifacts into a structured Opportunity Brief. Your goal is to [Aseparate signal from noise by grading incoming evidence against the RocketShip Strategy One-Pager, surfacing only strategy-aligned opportunities backed by real data.dd Main Goal].

Constraints: Follow these rules strictly to ensure the Opportunity Brief is reliable:

- Never invent metrics or quotes.
- Always show which source each insight came from (e.g., Artifact 4).
- If evidence for a claim is weak, say so explicitly.
- Never promote a request to ACTIONABLE unless it maps to at least one pillar of the RocketShip Strategy One-Pager with concrete supporting evidence (specific error codes, client names, or quantifiable metrics).

Few-Shot Example

Highlight the specific example you provided in the prompt.

Below is an example of the quality I expect. Use this as a reference for tone, density, and evidence-linking:

- **Example Problem:** Enterprise customers are experiencing recurring data pipeline failures, but the support tickets describing them are scattered across Slack, Jira, and call transcripts with no unified priority — causing a 48-hour average delay before the PM team even sees the pattern.
- **Example Evidence:** "Three Tier-1 enterprise accounts (Acme Corp, Orion Data, NovaSys) reported identical ETL timeout errors (Error Code: PIPE-4092) in the last 14 days. Support tickets #4501, #4523, and #4540 all reference the same connector version (v3.1.2). Call transcript from March 12 confirms Acme is evaluating a competitor due to unresolved downtime." — Sources: Artifact 2 (Support Ticket Log), Artifact 4 (Customer Call Transcript), Artifact 5 (Jira Backlog)
- **Example Risk:** We are assuming that Error Code PIPE-4092 is the root cause across all three accounts because the symptoms overlap, but no engineering diagnosis has confirmed a shared failure point. If these are actually three separate bugs, consolidating them into one PRD would under-scope the fix and delay resolution for two of the three clients. This is a Leap of Faith that requires engineering validation before the PRD moves to ACTIONABLE.

AI Strategy One-Pager

Link to Artifact from Module 2: ***Add your link here (make sure it's sharable!)***

https://docs.google.com/document/d/1cmYTpZD6oRwlWBsjkImFvmmfulnwEOg3gEqQ0BLC5VU/edit?tab=t.0

Completed Sections

- Problem and Workflow:
  - **The Problem:** Incoming signals from tickets, calls, and Slack pile up unclassified for days, stalling deals and trapping the PM team in reactive mode.
  - **Prevention:** Transform the multi-channel roar of transcripts, tickets, and calls into structured, evidence-backed clarity (system).
- Target Metrics:
  - **Resolution Cycle Time:** *Mean Time to Triage & Route (MTTR):* Reduce the average time from signal intake (ticket, transcript, or call) to classified, prioritized, and owner-assigned status from ~48 hours to under 60 seconds.
  - **Leadership Proof:** A 80% reduction in missed responses and 60% improved time to resolution.
- Autonomy Level:
  - **Autonomy Choice:** Copilot.
  - **Explicitly Avoiding:** Agent. No end to end autonomy. Agent will involve human in a loop based on complexity.
- Data & Model Approach:
  - **Approach:** RAG
  - **Explicit Avoiding:** We are avoiding LLM because of the lack of expertise in the company's topic, as it could hallucinate and lead to wrong decisions.
- Risks and Mitigations:
  - **Risk:** The model hallucinates, based on training data on past performed workflows and not prioritizing appropriately.
  - **Mitigation:** Use the "standardized workflow" document that will keep updating the use of the AI paired with a recurring weekly cadence of quality assurance.
- V1 Scope:

  IN: Prioritize work items and identify appropriate teams or resources to deliver responses.

  OUT: Will not do any communications outside of core deliverables and defined personas.

RAG Architecture & Data Strategy

Data Sources
List the specific sources Juno retrieves from in order to use RAG. Hint: Use the Data Requirements section of your AI PRD snippet from Module 3!

Sources: The system must access:

- Slack
- Notion
- Jira
- Support tickets
- the “standardized workflow” document

Quantity: Limit message ingestion to the specific artifact where the dispute occurred and the team member requesting the support to prevent cross-contamination of user history.

Data Strategy
Describe how Juno finds the right information to minimize hallucinations. Hint: Use the Model Requirements and AI User Experience sections of your AI PRD snippet from Module 3!

Juno grounded Trust Requirement for AI User Experience

Verification: The AI cannot just give a "Yes" or "No" answer. It must show the receipts, where every decision must point directly to the specific rule or photo it used to reach that conclusion.

Fail-Safe Behavior: If the AI can’t find a clear reason in the files to charge a guest, it must stop and say, "I don't have enough evidence to

UX Design & Interface

(Optional) Live Demo Link: *Add your link to Lovable or your UX design here (make sure it's sharable!)*
https://juno-insight-lab.lovable.app

Key **Screenshots**
(Paste 1-2 screenshots of Juno's UI) and your AI User Flow diagram from Module 4

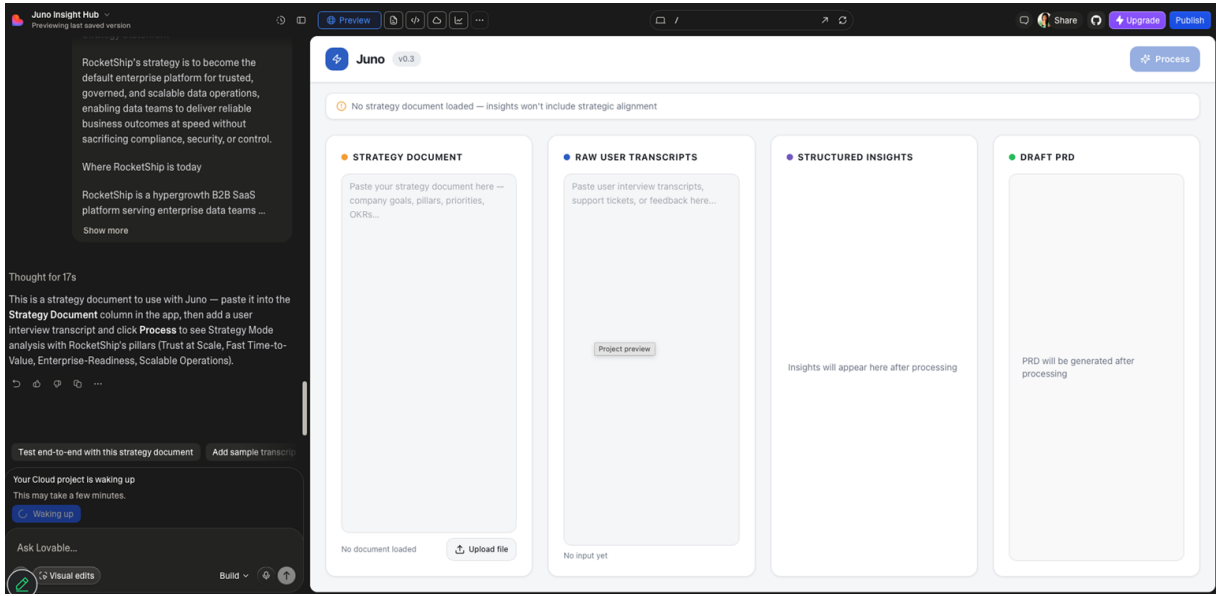

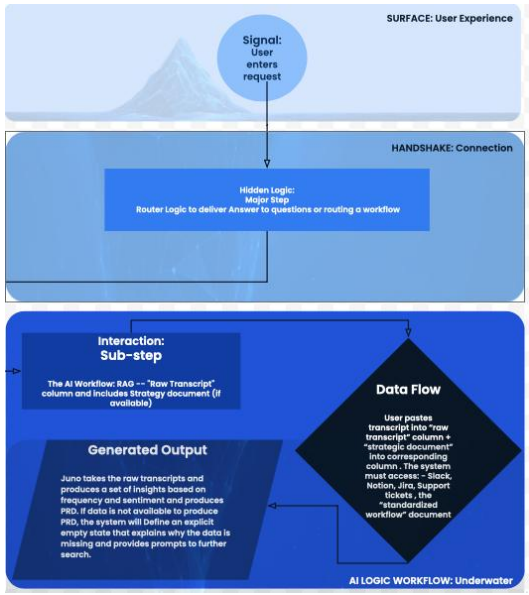

Agent Workflow Spec (AWSpec)

Link to Artifact: https://juno-insight-flow.lovable.app

https://docs.google.com/document/d/1x9vbPxEj6JQ4h9bMtc-npb-Lm2eLUvYHb5dZqmFF0GM/edit?tab=t.0

Completed Sections

Primary Agent Pattern: ReAct (Reason → Act → Observe → Re-plan)

Agent Roles:

- **Project Manager Agent (Juno):** Routes incoming requests, manages conversation flow, decides between strategy question mode and pipeline evidence mode
- **Strategy Owner Agent:** Audits inputs against the RocketShip Strategy One-Pager; grades each as ACTIONABLE, FUZZY, or VAGUE (P0–P3 priority)
- **PRD Writer Agent:** Generates formal PRDs only for ACTIONABLE items; refuses FUZZY/VAGUE requests
- **Human PM (L2 autonomy):** Final approval authority; provides missing info on FUZZY items

Key Steps:

1. **Reason** — Juno receives raw input and classifies it as a strategy question or pipeline evidence
2. **Act** — Strategy questions get answered directly from the strategy doc; pipeline items get routed to the Strategy Owner for grading
3. **Observe** — Green (ACTIONABLE) → PRD Writer drafts; Yellow (FUZZY) → asks PM for missing details; Red (VAGUE) → rejects with explanation
4. **Re-plan** — If FUZZY, PM provides clarification and Juno re-audits (max 2 loops, then escalates to manual discovery)

Evaluation Plan

Definition of “Good”

- *What are the 3 criteria Juno’s output must meet to be considered "Associate PM quality"?*
- *Mean Time to Triage & Route (MTTR):* Reduce the average time from signal intake (ticket, transcript, or call) to classified, prioritized, and owner-assigned status from ~48 hours to under 60 seconds.

Eval Methodology

- *How will you test Juno?*
- *Test: Review what Juno over-prioritizes, what it treats as equally important, and whether it distinguishes between superficial preferences and true workflow-blocking problems (Track via thumbs up-down and input of evaluation).*

Risk Mitigation

- *Describe one specific guardrail you would build to prevent Juno from making a "risky assumption" or hallucinating a technical blocker.*
- **Model hallucinating Mitigation: Using RAG**, Use the “standardized workflow” document that will keep updating the use of the AI paired with a recurring weekly cadence of quality assurance.

Group Insights: Building Process

Friction Points

- *Describe the biggest technical or design challenge you faced while building and designing Juno*
- ***Manging the tension between agentic capabilities and delivering a trusted governed output/solution***
- ***Maging the user transcription***
- ***Thinking through the feedback loop or validation process to ensure model wasn't learning from biased or noisy data.***

Key Learnings

- *What was your group's main "aha" moment during the project process?*
- ***The significant impact a prototyping approach can have on the End Users (Lovable tool alone was a "aha" moment)***
- ***The deliverable that can be derived from a solution like Juno to support a PM in quality product deliverables***
- ***The Aha was realizing that Juno wasn't just a front facing tool, but its ability to create meaningful PRDs and issue management.***

Aha! Moment

- *Share a few surprising discoveries or insight your group gained from the hands-on build labs.*
- ***A lot of discoveries during the first lab on the the code that is created and available within Juno to support building a prototype.***

Thank You