Juno AI User Workflow

Step-by-Step User Workflow

Signal: Surface Entry point(s)	• **Entry point 1:** Transcripts: User uploads all recent transcripts and tickets. • **Entry point 2:** Slack messages • **Entry point 3:** Email • **Entry point 4:** Telegram message, X mention, or Bounty Hunters feedback is ingested
Hidden Logic: Handshake Major step(s)	**Classify signals** and provides draft prioritisation output and weighting (hierarchy). 1. **Juno ingests fresh signals** from the connected sources: Slack threads (ability to select channels), all tickets (last 3 months), CRM / sales notes (active + lost deals), X mentions, Emails, Telegram, Bounty Hunters feedback 2. **Juno refreshes against a default hourly sync cycle,** with refresh frequency tweakable by source 3. **Juno retrieves the Top 3–5 most relevant segments** 4. Juno classifies signals into: • Revenue blocker • System risk • User friction 5. **Juno compares the evidence against the strategic business objectives** underwater: • Prevent opinion-driven prioritization • Weight high-impact issues over noise • Support the target outcomes: ○ Reduce monthly P0 escalations ○ Decrease critical bug resolution time 6. **Juno generates:** • Priority score • Ranked recommendation • Evidence-backed rationale • Rank confidence

Juno AI User Workflow

Interaction Maneuver: Surface Sub-step(s)	• **"Syncing latest signals..."**
	• **"Scanning sources..."**
	• **"Cross-checking strategy fit..."**
	• **"Scoring priority..."**
	• Lovable dashboard shows:
	  ○ Priority score
	  ○ Priority type badge: Revenue / System / User
	  ○ Rank confidence label
	  ○ Evidence snippets with source citations
	  ○ Recommended next action
	  ○ Conflict flag if sources disagree

Feedback Loop: Underwater Data Flow	Data flows in from the selected internal and external sources into Juno's evidence engine. Index is refreshed on a **default hourly cadence**, adjustable by source Retrieval narrows to the **Top 3–5 relevant segments** Juno cross-references retrieved evidence with prioritization logic from the One-Pager. The system stores: • Retrieved evidence • Scoring inputs • Confidence label • Final recommendation • PM override action, if any

Juno AI User Workflow

Feedback Loop: Underwater AI-Generated Output

- Defensible prioritization score
- Opportunity Brief summary
- Ranked recommendation
- Evidence table with sources/snippets
- Rank confidence
- “Insufficient evidence to prioritize—requires human validation.” when evidence is weak.

Note! Juno uses a Hybrid approach (Agentic + RAG): **RAG grounds insights** in real artifacts; **agentic handles multi-step reasoning**, cross-source validation, and conflict detection.

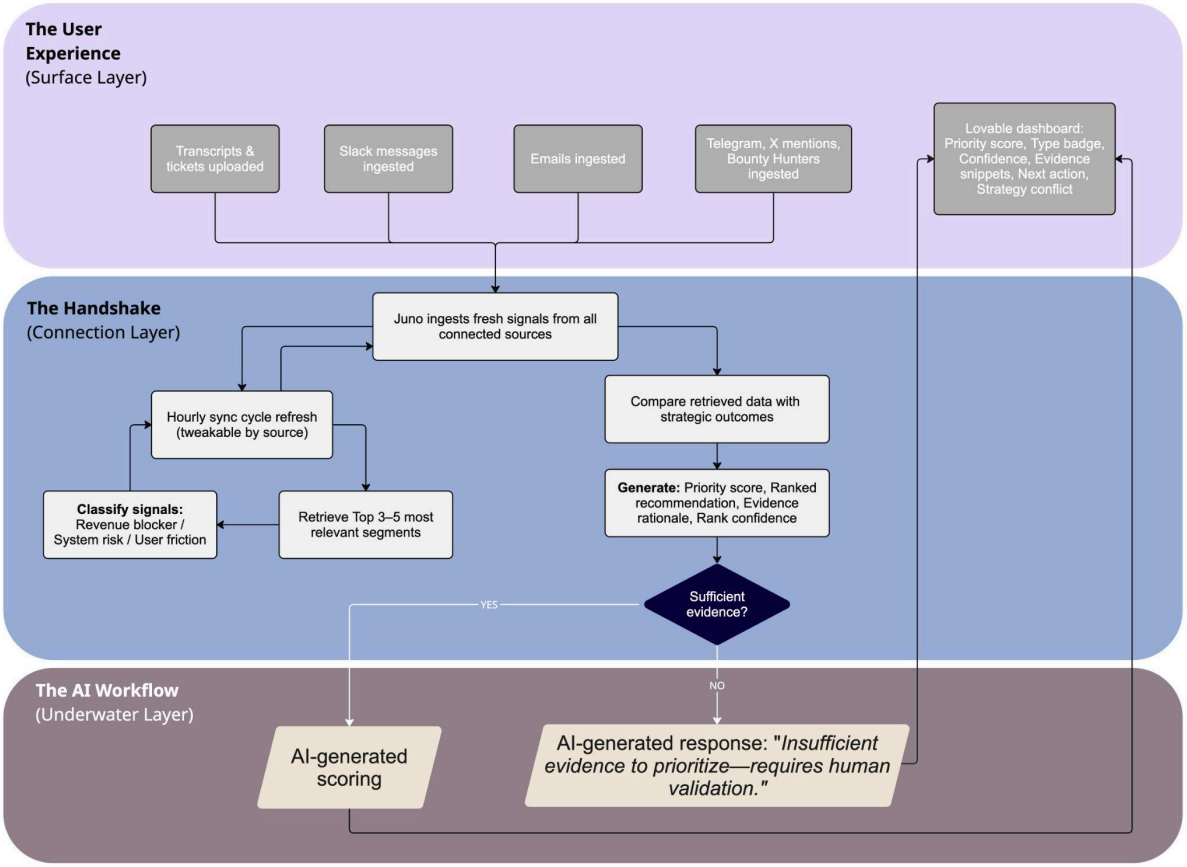

Juno AI User Workflow

Juno Workflow Specs

Field	Agent Spec

Actors	**For Ingestion Agent:**
	• Role:
	Fully-autonomous (L3) Ingestion agent that:
	  ○ Ingests signals
	• Responsibility:
	  ○ Evidence gathering
	
	For Synthesizes Agent:
	• Role:
	Fully-autonomous (L3) Synthesizes agent that:
	  ○ Synthesizes evidence
	• Responsibility:
	  ○ Multi-step reasoning
	  ○ Conflict detection
	
	For Prioritization Agent:
	• Role:
	Semi-autonomous (L2) prioritization agent that:
	  ○ Generates priority score + Opportunity Brief
	• Responsibility:
	  ○ Conflict detection
	  ○ Recommendation generation
	
	For PM (Human):
	• Role:
	  ○ Final decision-maker and validator
	  Responsibility (L2):
	  • Review Juno's output
	  • Override or approve prioritization
	  • Handle ambiguous / low-confidence cases

Juno AI User Workflow

Pattern Plan

Chosen Pattern: ReAct (Distributed Across Agents)

Instead of one loop, the ReAct pattern is distributed across agents:

Flow Across Agents

Step 1 — Ingestion Agent (Act)

- Triggered by:
  - P0 ticket
  - Slack escalation
  - CRM update
  - Email / X / Telegram / Bounty Hunters
- Actions:
  - Pull data from sources
  - Normalize + index
  - Pass structured signals forward

Step 2 — Synthesizer Agent (Reason → Act → Observe)

Reason

- Identify missing context
- Determine relevant sources

Act

- Retrieve Top 3–5 segments from:
  - Slack (selected channels)
  - Tickets (last 3 months)
  - CRM
  - External signals

Observe

- Classify signals:
  - Revenue blocker
  - System risk

Juno AI User Workflow

	◦ User friction • Detect conflicts: ◦ e.g., CEO request vs system fragility **Step 3 — Prioritization Agent (Reason)** • Apply strategy: ◦ Revenue > System > User • Align with business goals: ◦ Reduce P0 escalations ◦ Reduce bug resolution time **Step 4 — Output** Happy Path • Priority score • Ranked recommendation • Opportunity Brief • Evidence + snippets • Rank confidence Fail-Safe Path If weak evidence: → "Insufficient evidence to prioritize—requires human validation."
Memory	**Memory Types** **1. Contextual Memory (Per Request)** • Active signals being processed • Current prioritization session **2. Episodic Memory (Per Cycle)** • Current prioritization cycle (e.g., weekly roadmap)

Juno AI User Workflow

	3. Linked Memory (Shared Across Agents)
	
	• Live connections to:
	  ◦ Slack
	  ◦ CRM
	  ◦ Tickets
	  ◦ External sources
	
	Memory Controls (Critical)
	
	• Top-K Constraint: 3–5 segments only
	• Freshness Constraint: Hourly synced data (tweakable per source)
	• Confidence Threshold:
	  ◦ Low confidence → escalate to human
	• Stop Condition:
	  ◦ No strong evidence → no decision

Tools	• Ingestion Agent:
	  ◦ Slack API → Read-Only
	  ◦ CRM API → Read-Only
	  ◦ Ticketing System → Read-Only
	  ◦ Email / X / Telegram APIs → Read-Only
	• Synthesizer Agent:
	  ◦ Ingestion Agent -> Read-Only
	• Prioritization Agent:
	  ◦ Synthesizer Agent -> Read-Only
	  ◦ Juno UI (Dashboard) -> Write-Only