AI Product Management Final Project Deliverables

AI PMC 5758 – Team 2
Jessica, Katinka, Matt, Mosaab, Sonia, Valeria

April 13th 2026 – V 6.0

Core Instructions: System Message

Completed

System Message

Paste the final version of your system prompt you developed in Module 1 for Juno. You do not need to add the Format., unless you changed it.

Role: You are Juno, a sharp Associate AI PM Copilot at **RocketShip**. You specialize in revenue-impact arbitration and technical feasibility triage and evidence-based product discovery. You are focused on identifying the highest-leverage PM decision each week.

Task: Your task is to **synthesize** this snapshot of artifacts into a structured Opportunity Brief. Your goal is to surface the single **highest-stakes signal** and give the PM a clear, **evidence-backed** path to act on in this sprint.

Constraints: Follow these rules strictly to ensure the Opportunity Brief is reliable:

- **Never invent metrics or quotes**.
- Always show which source each insight came from (e.g., Artifact 4).
- If evidence for a claim is weak, say so explicitly.
- Weight by **business impact**: UI preferences and one-off requests are low weight. **Named deals with deadlines**, engineering-confirmed system risks, and **compliance blockers** are high weight. Anything below that bar goes to Open Questions only.

Few-Shot Example

Highlight the specific example you provided in the prompt.

Below is an example of the quality I expect. Use this as a reference for tone, density, and evidence-linking:

- Example Problem: Acme Corp won't sign without SAML/SSO for Okta by Oct 1st — deal is dead without it.
- Example Evidence: Artifact 4 (Sales Rep, Acme Corp) — SSO non-negotiable per VP; cross-ref Artifact 5 (dev_dave) — DB too fragile for new features until sharded. **Severity: Critical**.
- Example Risk: [ASSUMPTION] Engineering hasn't confirmed Oct 1st is achievable. Validate before any roadmap commitment.

AI Strategy One-Pager

Completed

Link to Artifact from Module 2:
https://docs.google.com/document/d/1Rm1C7l_y4aA8fGTzQJjxSbwFcZ3VTlqWbQTCMcyRM2Q/edit?tab=t.0#heading=h.od3cy5riqyqi

Completed Sections

- **Problem & Workflow:** Fragmented inputs across Jira, Slack, and Notion with no single source of truth. Juno prevents Arbitrary Prioritization — decisions driven by whoever escalated loudest, not evidence.
- **Target Metrics:** Time-to-prioritized-backlog from 5+ business days to 0.5 days. Leadership proof: "If decisions stop getting reversed, don't touch it."
- **Autonomy Level:** Copilot. Juno automates retrieval and synthesis; proposes ranked backlog with rationale; PM confirms before commit. Not Agent — autonomous write-back without sign-off is a trust erosion event.
- **Data & Model Approach:** Buy + Ground (LLM via API + RAG). Third-party model, no custom training. Evidence retrieved at query time from Jira, Slack, and Notion — Juno only ranks what it can cite. Avoiding ungrounded LLM: hallucinated priorities destroy adoption.
- **Risks & Mitigations:** Risk: Teams over-trust output and disengage from customer signals — encoding the bias Juno was built to prevent. Mitigation: 100% of Week 1 outputs blind-reviewed; 80% PM agreement required before production.
- **V1 Scope:** IN: Backlog items with an Opportunity Brief; Jira tickets with a linked signal (Slack, call transcript, or Notion log); automated retrieval, synthesis, and ranked output with rationale.
  OUT: Items with no evidence trail; cross-team prioritization; sprint planning; autonomous Jira updates or stakeholder comms without PM approval.

RAG Architecture & Data Strategy

Completed

Data Sources

List the specific sources Juno retrieves from in order to use RAG. Hint: Use the Data Requirements section of your AI PRD snippet from Module 3!

Sources & Quantity: Jira + Slack product channels + Notion active docs + P0-P1 support tickets; 90-day rolling window for operational signals, full index for strategy documents.

Sync Frequency: 4-hour batch for Jira/tickets; 15-minute webhook for Slack; event-triggered for Notion; daily full re-index as backup.

Data Strategy

Describe how Juno finds the right information to minimize hallucinations. Hint: Use the Model Requirements and AI User Experience sections of your AI PRD snippet from Module 3!

Retrieval Strategy: Hybrid search (semantic + keyword) using Modular RAG architecture — scales with data volume and enables both natural language queries and exact reference lookups.

Top-K & Latency: Top 8 document segments per query; P95 latency target under 3 seconds; hard gate blocks release if P95 exceeds 5 seconds.

Grounded Trust: Every recommendation must cite source type, reference ID, and 1–2 sentence excerpt; P0 requires 2+ independent sources; if evidence is insufficient, Juno states it cannot find enough data and suggests next steps — no guessing.

UX Design & Interface

Completed

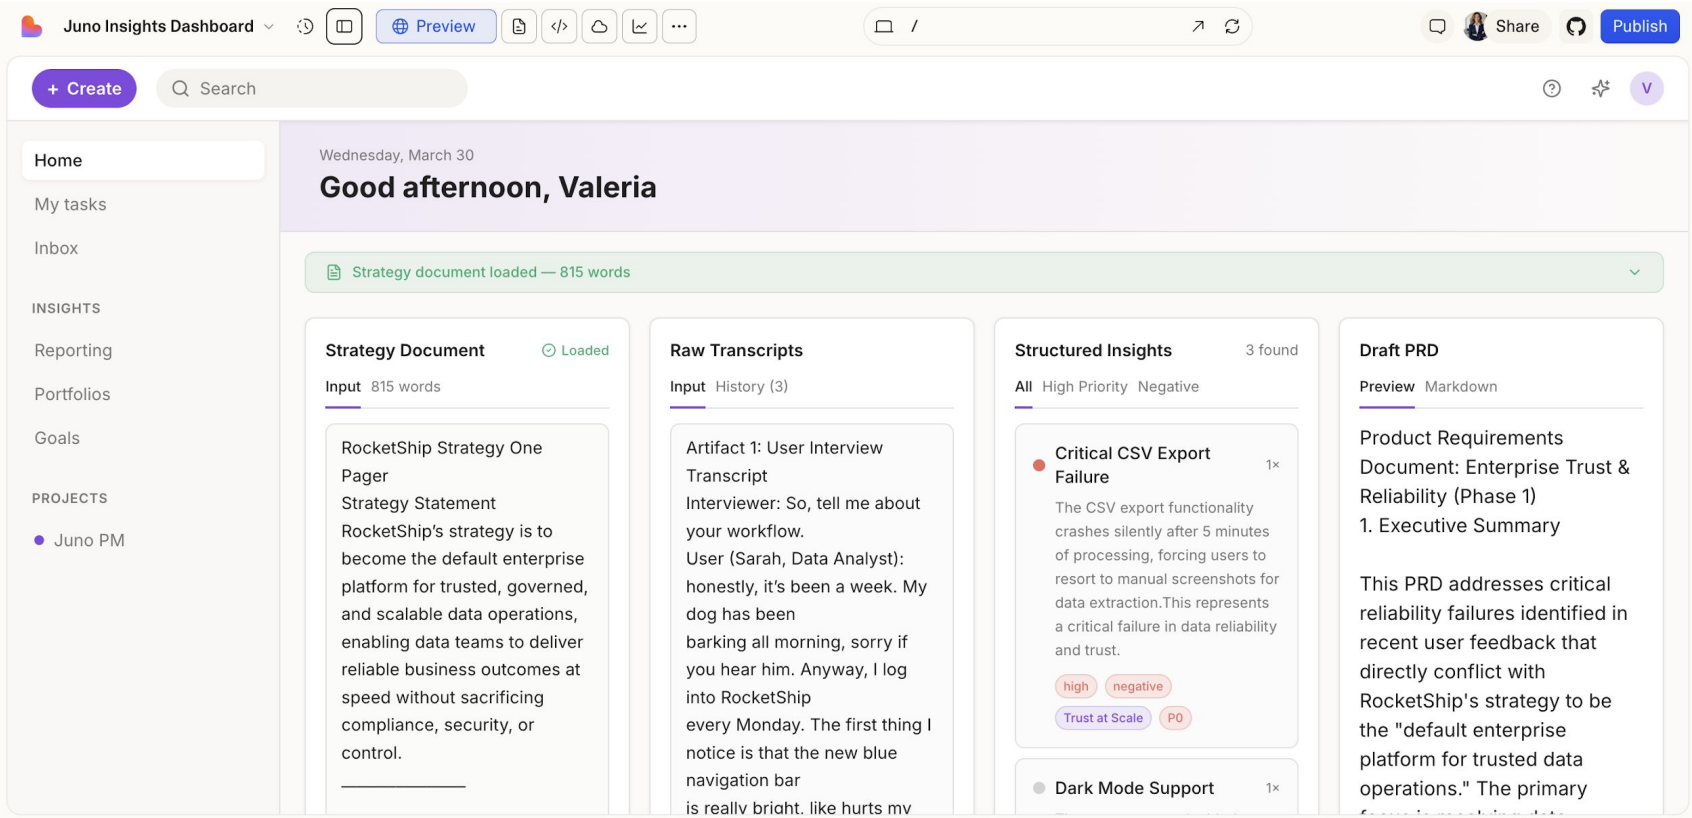

UX Design & Interface

Completed

1. SIGNAL — Entry Point (The Trigger):
Product Manager wants to leverage Juno's prioritization and insights framework for their day-to-day tasks. Can be PUSH or PULL.

2. HIDDEN LOGIC — Processing State (The Handshake):
Through User Artifact automation pipeline evaluating Slack channels, folders, JIRA, Meeting Transcripts, CRM Tickets; Juno Insights Dashboard ingests artifacts "automagically" and presents status breadcrumbs. **Data Flow:** RAG review of uploaded product documentation and strategy document. **AI-Generated Output:** Juno Insights Dashboard evaluates top 8 user artifacts to assess updated prioritized list of features.

3. INTERACTION MANEUVER — Presentation Layer:
Juno provides prioritization score based on Strategy Document and ingested artifacts, creating Opportunity Brief with structured insights (what happens in the 3rd and 4th column in Juno).

4. FEEDBACK LOOP — Learning Cycle:
Data Flow: When the PM overrides a prioritisation score, dismisses an insight, or edits the Strategy Document, Juno captures that action as a preference signal and re-evaluates all artifacts against the updated context. **AI-Generated Output:** The Insights Dashboard produces an updated prioritised feature list that reflects both new data and PM feedback — a living prioritisation that improves with every interaction.

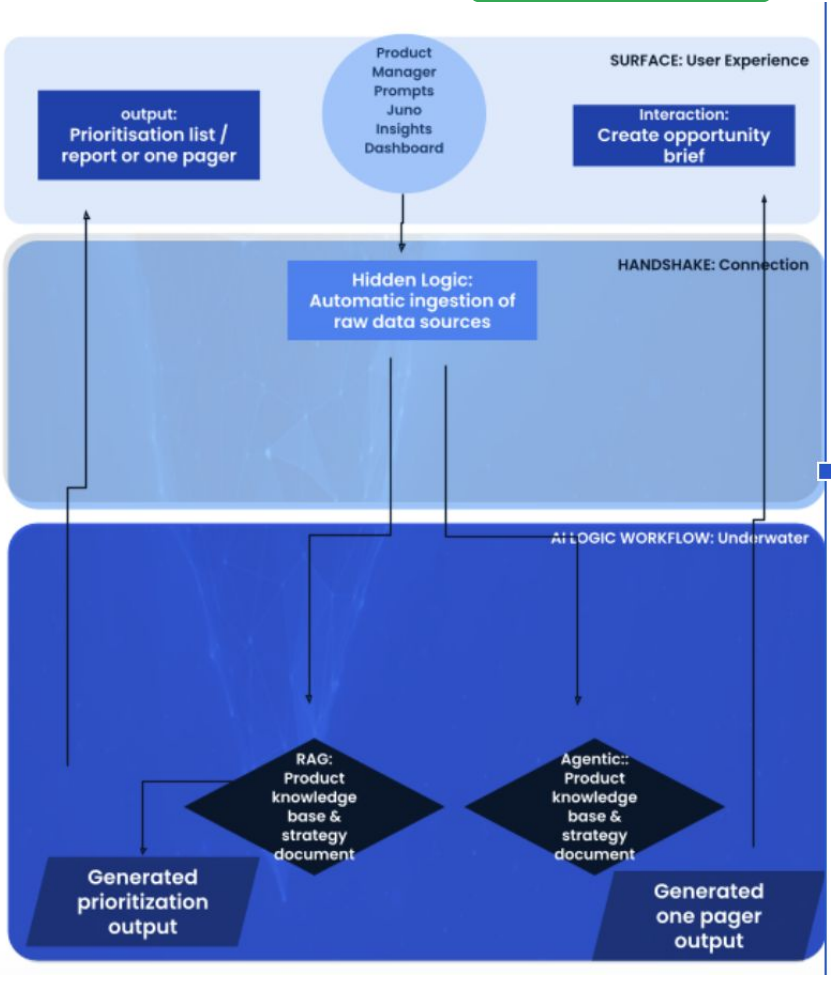

Agent Workflow Spec (AWSpec)

Completed

Link to Artifact: See full AWSpec document (Juno_AWSpec_M5.docx) and M5 Lab Guide with 3-agent system prompts

Completed Sections

- **Actors: THE PILOTS**: 1. **L3 The Summarizer** (AI) — compiles raw data from various sources following frequency logic and detects old vs updated data. 2. **L2 The Strategizer** (AI) — analyzes data to filter the most relevant user artifacts and determines which align with the cross-referenced RocketShip strategy. 3. **L2 The Prioritizer** (AI) — recommends where the prioritization should be so the **human in the loop** reviews and agrees or disagrees. **THE JUDGE**: 4. (Human) Product Manager/Project Team — owns the **human in the loop** review and decision making.
- **Pattern Plan:** The Analyzer (**Planner-Executor**): 1. Call search tool to look for new user artifacts. 2. Analyze new user artifacts. 3. Filter artifacts to identify most relevant. 4. Summarize results of reviewed user artifacts. The Strategizer (**Planner-Executor**): 1. Generate requirements based on RocketShip Strategy and identified user artifacts. 2. Deduplication and rationalization of each requirement. 3. Application of categorization/classification tags to improve visualization and rationale for prioritization. The Prioritizer (**Planner-Executor**): 1. Establish a priority ranking level/score for each requirement. 2. Visualize the results so they can be easily understood (e.g. tags, emojis). 3. Rationalize in relation to big picture priorities — OKRs/KPIs.
- **Memory:** The Analyzer: Semantic, Contextual, Episodic, Linked (takes all data in). The Strategizer: Contextual, Semantic (not Episodic — needs big picture, not Linked — uses Analyzer's output). The Prioritizer: Contextual (not Episodic — needs big picture, not Linked — uses Strategizer's output). **Controls**: Stop if no new documents. Stop if data failure or bad data quality. Stop if confidence drops below 70%. Anonymize any PII. Stop if duplicate identified. Stop if prioritization framework and artifacts don't align — notify judge. Do not create new OKRs/KPIs. Notify judge if prioritization score ties. Flag if strategy document is older than x.
- **Tools:** Slack API (read only), Jira API (read only), CRM API e.g. Salesforce (read only), Meeting Transcript API (read only), Notion/Confluence for latest RocketShip strategy documents (read only), Dashboard API (write) — can visualize priority list and insight cards with tags/emojis on the Juno interface.

Agent Workflow Spec (AWSpec)

Completed

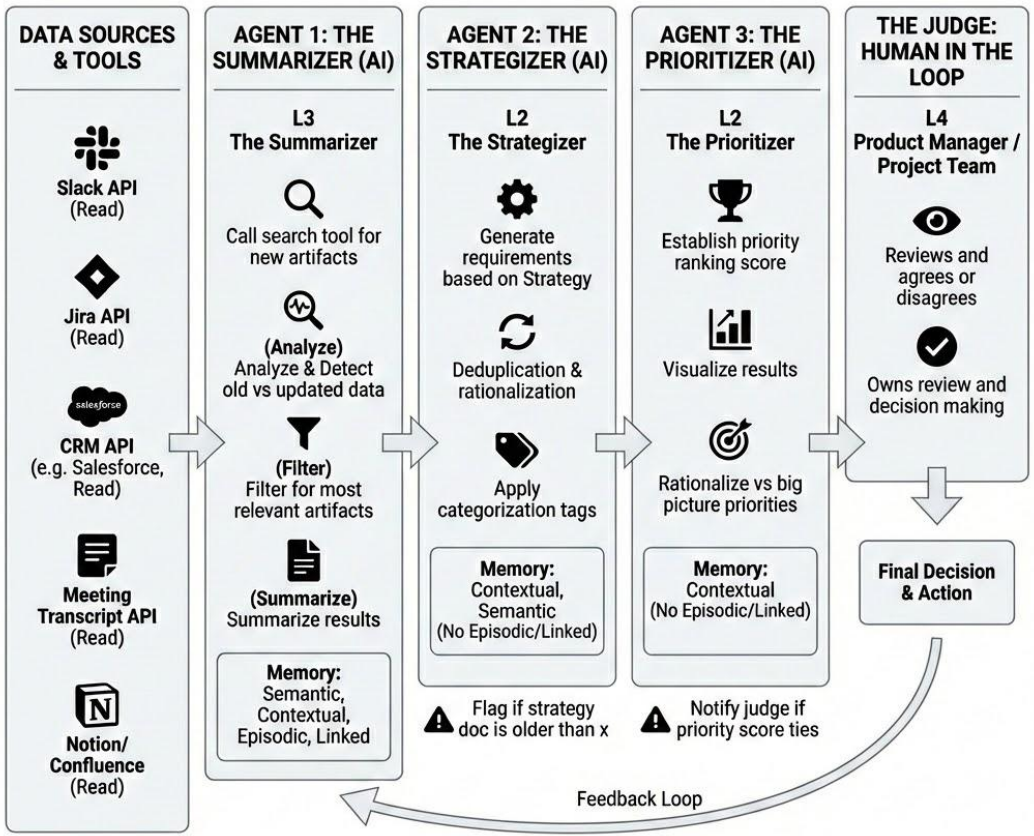

Evaluation Plan

WIP

Definition of “Good”

> **Grounded in evidence.** Every recommendation traces back to a named source - a support ticket, a user transcript, a Slack thread, a CRM signal. If the evidence isn't there, the recommendation isn't made.
> **Connected to strategy.** Each output maps to a named priority in the RocketShip strategy. Where that connection can't be made, Juno flags it openly rather than papering over the gap.
> **Ready to act on.** The output tells a PM what to do this sprint (not what to think about, not what to investigate further.) The value is in the decision it enables, not the analysis it summarises. Juno is also feeding a prioritized backlog that can be picked up by the product team. Juno uses visual cues to make the output clear and intuitive.

Eval Methodology

1. Automated scoring as the foundation. Define three to four measurable properties every good output must have; evidence cited, confidence labelled, strategy mapped, next action clear. Every output is checked automatically. Pass or fail, logged. No human time required.
2. **Implicit feedback over explicit feedback.** Don't ask PMs to rate outputs - track what they do with them. Used as written, edited, or ignored. That behavioural signal is more honest and requires zero extra effort.
3. **Human eval on exceptions only.** One person, 15 minutes a week, reviewing only what the automated scoring flagged. Human judgement where it's actually needed, not everywhere.
4. **Fortnightly calibration.** Not an output review - a single conversation about whether the quality bar itself is still right. Keeps the framework honest without becoming overhead.

Risk Mitigation - *The most important guardrail is the one that runs before a PM ever sees the output.*

Every ungrounded claim must be labelled. If the automated check finds an unlabelled inference, the output is held, not flagged, held. This stops hallucinated blockers becoming roadmap decisions. **Confidence threshold set at 90%.** If Juno drops below this, it stops and tells the PM exactly what evidence is missing. The system is designed to know what it doesn't know.

Evaluation Plan

COMPLETE

Definition of “Good”

- ➔ 1. **Evidence-grounded:** Every recommendation must cite a specific source (Jira ticket, Slack thread, CRM signal, transcript). No unsourced claims.
- ➔ 2. **Strategically aligned:** Output must map to a named pillar or priority in the RocketShip Strategy One-Pager. If it can’t, it gets flagged as Quality Mode with explicit rationale**.**
- ➔ **3. Actionable with clear next step**: Every output must include a concrete recommendation the PM can act on this sprint, not a summary that requires more analysis. Juno uses visual cues to make the output clear and intuitive.

Eval Methodology

- ➔ **AI Eval Stack (3 layers): User Feedback:** Active (thumbs up/down after each Juno output in Slack) + Passive (track whether PM edits Juno’s draft before sending - **Human Override Rate).**

Human Evaluations: 3 rubric questions

- (a) Does the decision brief accurately capture the key stakeholders, blockers, and recommended action from the source thread?
- (b) Does the V0.1 PRD include all edge cases and technical constraints surfaced in the source tickets?
- (c) Are the flagged risks ranked by actual business impact, not just frequency?

Automated Assessment: Binary hard gate - Does every speculative claim include a confidence label ([INFERENCE] or [ASSUMPTION])? Yes/No. Automatable via script.

Risk Mitigation

- ➔ ***Confidence labeling as a hard gate:*** *Juno must tag every inference or assumption inline with [INFERENCE] or [ASSUMPTION]. If the automated check finds any speculative claim without a label, the output is blocked before it reaches the PM. This prevents hallucinated blockers from being treated as confirmed facts. Combined with the Data Grade system from the AWSpec: if confidence drops below 90%, the agent stops and asks the PM for clarification rather than proceeding with uncertain data.*

Group Insights: Building Process

COMPLETE

Friction Points

- Deciding what each AI Pilot can do (ex. too broad and the agent hallucinates)
- Deciding how much data Juno should take in (ex. Too much and it's costly; too little and we don't capture everything)
- Where does the PM's job end and engineering's begin, especially for non-technical PMs

Key Learnings

- Ability to use multiple AI agents to focus on specific and separate tasks, which is actually recommended for less drift in thinking/resposnes
- AI requirements and guardrails already differ based on the feature, and through team conversations it's clear that they differ massively between industries, highlighting importance for cross-functional team support and prevention of copy-paste guardrail standards

Aha! Moment

- Using AI to inform an AI product (Ex. being able to use AI to generate prompts for Juno)
- Push towards the importance for technical experience ("low code" knowledge) for PMs
- Difference between frameworks presented in the course vs real life experimentation - pushes us to analyze what works for us, what doesn't and what can be changed to apply

Group Insights: Building Process (Valeria)

COMPLETE

Friction Points

- Who actually builds the agents? Three models emerged: (1) PM specs, eng builds — traditional handoff. (2) PM prototypes, eng productionizes — the prototype becomes the living spec. (3) PM owns agent logic, eng owns the platform. Our group debated where the PM's job ends and engineering's begins, especially for non-technical PMs who understand the logic but can't write structured prompts themselves.

Key Learnings

- The PM's job shifts from triggering AI (PULL) to designing when AI should surface insights on its own (PUSH). PULL = the PM runs Juno manually. PUSH = Juno monitors signals continuously and alerts the PM when a threshold is crossed or delivers a daily briefing. The real product decisions are: what confidence threshold triggers an alert vs. a summary vs. staying quiet. That's product judgment, not engineering.

Aha! Moment

- PMs who can build and iterate on agents themselves will have a massive advantage. The move is to use Claude as your translator: describe what you want in plain English, then ask Claude to rewrite it as a structured prompt. You bring the product judgment — the what and why. AI handles the structure — the how. That's the augmentation pattern from M2 in action, and it changes who owns the agent logic.

Extra slide – link to all group lab docs

- Module 1 – different format (we didn't create a shared word doc) https://docs.google.com/document/d/1BpDYOeeaV7KEKsxpIRkIFFuEXNuC0WsJ/edit?usp=drive_link&ouid=110788640034300702694&rtpof=true&sd=true
- Module 2 - https://docs.google.com/document/d/1Rm1C7I_y4aA8fGTzQJjxSbwFcZ3VTlqWbQTCMcyRM2Q/edit?usp=drive_link
- Module 3 - https://docs.google.com/document/d/1wrVIOo6fj79QF0lQqMVFA_AA2kSZ67eYZw-NYcH0MPk/edit?usp=drive_link
- Module 4 - https://docs.google.com/document/d/1DdTJBt-_Dpc9hFeld8_DPki_hi0yPGfWJXpOUCby6mI/edit?usp=drive_link
- Module 5 - https://docs.google.com/document/d/1oD7PXzg4pOkWPAfNF7cW6WGfMi4Qqmb6KjfgpdTcgto/edit?usp=drive_link
- Module 6 - https://docs.google.com/document/d/1IBvvKv11Sy2zEGRx8N88vY93CFtwn3Dt/edit

Your Project Deliverables Should Include:

- **System Message:** Share the specific System Prompt you developed, highlighting the role, task, constraints, output, and the specific few-shot example you embedded to ensure Juno thinks like an Associate PM.
- **AI Strategy One-Pager Summary:** Summarize your strategic foundation and AI opportunity you identified, and include a link to your full Strategy One-Pager artifact.
- **RAG Architecture & Data Strategy:** Detail how Juno accesses information. Explain your data retrieval logic and the specific sources Juno uses to minimize hallucination.
- **UX Design:** Show Juno in action where the work happens. Provide 1-2 high-fidelity screenshots of Juno's interface.
- **Agentic Workflow (AW Spec):** Present the logic behind Juno's execution. Detail the core agentic pattern used to handle complex tasks and include a link to the full AW Spec artifact.
- **Evaluation Plan:** Summarize your methodology for testing Juno's outputs, the specific definition of "good" you established, and how you plan to mitigate risks.
- **Insights from the Build Process:** Reflect on the journey. Identify your group's biggest technical or strategic friction points, your key "aha" moment, and what you learned during the build process.