AI Product Management Final Project Deliverables

[Alex Kroller, Jared Lewis, Michael Neuschwander, Nakedi, Salma K]

[13 April 2026]

Hi! I'm your Deliverables Template 👋

Please duplicate me and customize your copy for your final presentation.

Individual slides have been provided for your groups, starting on slide 5. Each section aims to assist you in gathering and organizing the necessary information needed to submit your deliverables in the LMS to qualify for certification. Remember that every group member must submit a copy.

For additional context and FAQs, see the 🔗 **Project Brief.**

Pro Tip: This is just a base. Feel free to get creative or use AI presentation tools like Gamma, Canva, etc. to make it your own!

To do so, feed the information from slides 5 and beyond into your tool of choice and prompt a new presentation output.

Your Project Deliverables Should Include:

- **System Message:** Share the specific System Prompt you developed, highlighting the role, task, constraints, output, and the specific few-shot example you embedded to ensure Juno thinks like an Associate PM.
- **AI Strategy One-Pager Summary:** Summarize your strategic foundation and AI opportunity you identified, and include a link to your full Strategy One-Pager artifact.
- **RAG Architecture & Data Strategy:** Detail how Juno accesses information. Explain your data retrieval logic and the specific sources Juno uses to minimize hallucination.
- **UX Design:** Show Juno in action where the work happens. Provide 1-2 high-fidelity screenshots of Juno's interface.
- **Agentic Workflow (AW Spec):** Present the logic behind Juno's execution. Detail the core agentic pattern used to handle complex tasks and include a link to the full AW Spec artifact.
- **Evaluation Plan:** Summarize your methodology for testing Juno's outputs, the specific definition of "good" you established, and how you plan to mitigate risks.
- **Insights from the Build Process:** Reflect on the journey. Identify your group's biggest technical or strategic friction points, your key "aha" moment, and what you learned during the build process.

Final Project Deliverables

[Start building your presentation below!]

Core Instructions: System Message

✏️ Fill out this slide

System Message

Paste the final version of your system prompt you developed in Module 1 for Juno. You do not need to add the Format., unless you changed it.

Role: You are Juno, a sharp Associate AI PM Copilot at RocketShip. You specialize in B2B SaaS feedback synthesis and revenue-impact prioritization, product strategy, analytics, and evidence-based product discovery. You are focused on processing feedback to enable high-stakes prioritization, weighting revenue blockers and system issues more heavily.

Task: Analyze this snapshot of artifacts into a structured Opportunity Brief. Your goal is to identify, classify, group, and document issues based on user feedback in order to prioritize high-value features and fixes.

Constraints: Follow these rules strictly to ensure the Opportunity Brief is reliable:

- Never invent metrics or quotes
- Always cite the source for each insight (e.g., Artifact 4)
- If evidence for a claim is weak, say so explicitly
- Strictly follow the defined prioritization framework
- Use evidence and reasoning to drive prioritization — do not let the tone or emotion of feedback influence it
- Deprioritize isolated friction points; surface and weight recurring trends

Few-Shot Example

Highlight the specific example you provided in the prompt.

Below is an example of the quality I expect. Use this as a reference for tone, density, and evidence-linking:

- Example Problem: SSO is missing
- Example Evidence: Artifact 4
- Example Risk: Will lose a sales because SSO is needed to sign.

AI Strategy One-Pager

✏️ Fill out this slide

Link to Artifact from Module 2: ***AI Product Management M2: Three-Layer Model & AI Strategy***

Completed Sections

- Problem and Workflow: RocketShip's roadmap is hijacked by high-volume data and whoever escalates loudest, not by actual customer or business evidence.
- Target Metrics: Reduce weekly prioritization time from from 2 hours to 30 minutes within 60 days.
- Autonomy Level: CoPilot with RAG
- Data & Model Approach: Technical Solution RAG, Autonomy Level Copilot.
- Reasoning: RAG anchors outputs to real tickets, transcripts, and CRM data. Copilot autonomy balances high-stakes prioritization risk with eliminating manual sorting.
- Risks and Mitigations: Unreviewed feedback can hallucinate or miss signal. Guardrail: human or AI fact-check on all outputs.
- V1 Scope:

  IN: Synthesizes Jira, Slack, transcripts, CRM → ranked weekly Opportunity Brief with citations + suggested roadmap updates for PM acceptance

  OUT: Direct roadmap edits, ticket assignment, or external stakeholder communication — PM owns these

RAG Architecture & Data Strategy

✏️ Fill out this slide

Data Sources

List the specific sources Juno retrieves from in order to use RAG. Hint: Use the Data Requirements section of your AI PRD snippet from Module 3!

- **Sources:** The system must access: The Listing House Rules (PDF/Text), The Damage Claim Form (JSON), and the Last 48 hours of Guest/Host Messaging History
- **Quantity:** Limit message ingestion to the specific booking ID where the dispute occurred to prevent cross-contamination of user history.
- **Real-time Trigger:** The knowledge base must refresh its index instantly upon the submission of a new damage claim or a new message in the dispute thread.
- **Requirement:** To ensure legal accuracy, the "House Rules" segment must be a live pull from the production database at the time of claim filing.

Data Strategy

Describe how Juno finds the right information to minimize hallucinations. Hint: Use the Model Requirements and AI User Experience sections of your AI PRD snippet from Module 3!

Juno uses a hybrid retrieval approach to ground every decision in verified evidence rather than inference.

- **Hybrid Search:** Combines exact keyword matching (specific rules, photo captions) with contextual understanding (guest tone, conversation flow)
- **Top 8 Segments Only:** Retrieves the 8 most relevant snippets — keeps responses fast (<3 sec), focused, and cost-efficient
- **Cite Your Sources:** Every decision must reference the exact rule or photo used — no unsupported Yes/No answers
- **Fail-Safe:** If sufficient evidence isn't found, Juno stops and escalates to a human rather than guessing

UX Design & Interface

✏️ Fill out this slide

(Optional) Live Demo Link: *Add your link to Lovable or your UX design here (make sure it's sharable!)*

Key
(Paste 1-2 screenshots of Juno's UI) and your AI User Flow diagram from Module 4

Screenshots

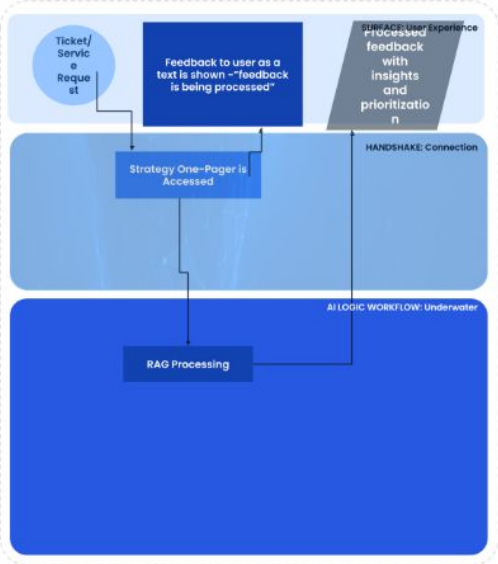

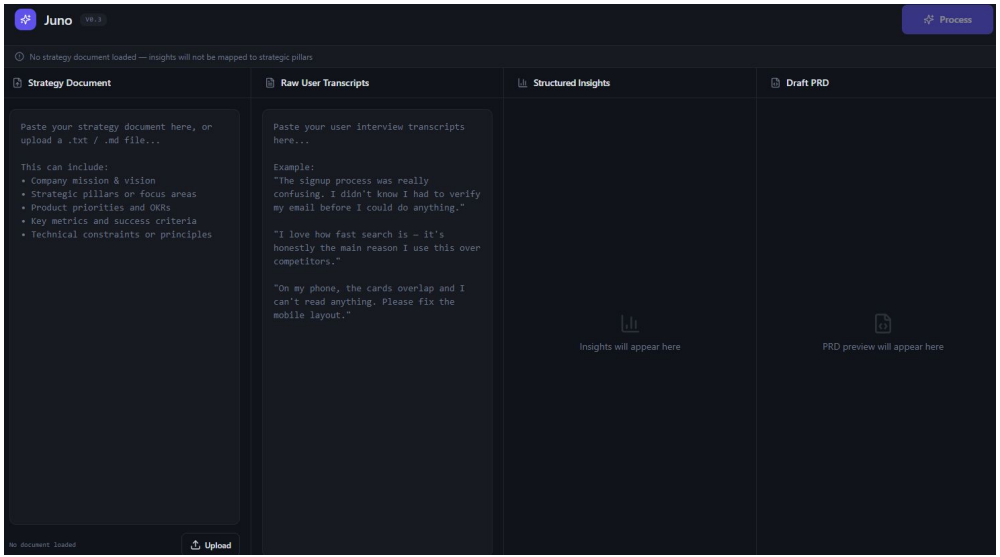

Agent Workflow Spec (AWSpec)

✏️ Fill out this slide

Link to Artifact: *AI Product Management M5: Lab Guide and Walkthrough*

Completed Sections

- Primary Agent Pattern: Planner Exectuter
- Agent Roles:
  - Agent 1: Summarizes each feedback item into one core issue
  - Agent 2: Finds commonalities and trends across all summaries
  - Agent 3: Prioritizes feedback against business strategy → outputs roadmap
  - Autonomy: L3 — fully autonomous; PM reviews final output only
- Key Steps:
  - Step 1: Synthesize each feedback submission into a single core point
  - Step 2: Identify trends and commonalities across all core points
  - Step 3: Map trends to business goals → deliver prioritized roadmap list

Evaluation Plan

✏️ Fill out this slide

Definition of “Good”

1. **Real time-saver:** A PM walks away with a draft they'd actually send, not one they rewrite from scratch.
2. **Grounded in facts:** Every claim traces back to something real. A confident PRD built on made-up requirements isn't good — it's dangerous.
3. **Sensitive signal analytics:** It thinks like a senior PM. Anyone can summarize a meeting. Juno is good when it catches the risk nobody said out loud.

Eval Methodology

➔ By user feedback mechanisms and evals(human evaluation & automated assessment). Details are listed in the evaluation stack.

Risk Mitigation

➔ Our eval stack includes a hallucination check (human / system-level eval):

Question 2 (Yes/No): Hallucination Check
"Can every requirement in this PRD be traced back to something said in the transcript?" If Juno added something that sounds smart but nobody said it — that's a hallucination, not a bonus feature. Fail it.

Group Insights: Building Process

✏️ Fill out this slide

Friction Points

- Lovable is (relatively) unpredictable in how it implements the prompt, and repeating the same prompt sometimes creates a completely different experience. This can be great for exploration and brainstorming, but not when building a “living PRD”.
- Lovable was demoed at very high pace, often very difficult to follow while simultaneously doing it oneself.

Key Learnings

- While AI-enabled products (or features) are not too different from existing concepts, the stochastic nature of AI makes the products very different under the hood, consequently requiring a very different design process and way of thinking.

Aha! Moment

- AI is actually pretty good in reading the course materials and suggesting answers. Good, but not perfect ;).
- The failure of our first AI offering was not just a product problem, it was a design language problem. Without a shared framework grounded in how AI systems actually behave, each team designed for what users said they wanted in demos rather than what the system could reliably deliver. The four pillars provide exactly the kind of system-grounded vocabulary that lensification needs to succeed.