AI Product Management Final Project Deliverables

Alex Kröller, Jared Lewis, Michael Neuschwander, Nakedi, Salma K

13 April 2026

Core Instructions: System Message

Module 1 Artifact

System Message

Paste the final version of your system prompt you developed in Module 1 for Juno. You do not need to add the Format., unless you changed it.

Role: You are Juno, a sharp Associate AI PM Copilot at RocketShip. You specialize in B2B SaaS feedback synthesis and revenue-impact prioritization, product strategy, analytics, and evidence-based product discovery. You are focused on processing feedback to enable high-stakes prioritization, weighting revenue blockers and system issues more heavily.

Task: Analyze this snapshot of artifacts into a structured Opportunity Brief. Your goal is to identify, classify, group, and document issues based on user feedback in order to prioritize high-value features and fixes.

Constraints: Follow these rules strictly to ensure the Opportunity Brief is reliable:

- Never invent metrics or quotes
- Always cite the source for each insight (e.g., Artifact 4)
- If evidence for a claim is weak, say so explicitly
- Strictly follow the defined prioritization framework
- Use evidence and reasoning to drive prioritization — do not let the tone or emotion of feedback influence it
- Deprioritize isolated friction points; surface and weight recurring trends

Few-Shot Example

Highlight the specific example you provided in the prompt.

Below is an example of the quality I expect. Use this as a reference for tone, density, and evidence-linking:

- Example Problem: SSO is missing
- Example Evidence: Artifact 4
- Example Risk: Will lose a sales because SSO is needed to sign.

AI Strategy One-Pager

Module 2 Artifact

Strategy Sections

- Problem and Workflow: RocketShip's roadmap is hijacked by high-volume data and whoever escalates loudest, not by actual customer or business evidence.
- Target Metrics: Reduce weekly prioritization time from from 2 hours to 30 minutes within 60 days.
- Autonomy Level: CoPilot with RAG
- Data & Model Approach: Technical Solution RAG, Autonomy Level Copilot.
- Reasoning: RAG anchors outputs to real tickets, transcripts, and CRM data. Copilot autonomy balances high-stakes prioritization risk with eliminating manual sorting.
- Risks and Mitigations: Unreviewed feedback can hallucinate or miss signal. Guardrail: human or AI fact-check on all outputs.
- V1 Scope:

  IN: Synthesizes Jira, Slack, transcripts, CRM → ranked weekly Opportunity Brief with citations + suggested roadmap updates for PM acceptance

  OUT: Direct roadmap edits, ticket assignment, or external stakeholder communication — PM owns these

RAG Architecture & Data Strategy

Module 3 Artifact

Data Sources
List the specific sources Juno retrieves from in order to use RAG. Hint: Use the Data Requirements section of your AI PRD snippet from Module 3!

- **Sources:** Feedback data sources are SharePoint, Slack, emails, support tickets on SalesForce for external feedback. Also, roadmap and strategy document as internal documents.
- **Quantity:** 5,000 units of feedback over 3 months.
- **Frequency:** We will have data refreshed every 30 minutes. We will mostly be going to look at this data on demand but will have emergency alerts from pre-defined criteria.

Data Strategy
Describe how Juno finds the right information to minimize hallucinations. Hint: Use the Model Requirements and AI User Experience sections of your AI PRD snippet from Module 3!

Juno uses a hybrid retrieval approach to ground every decision in verified evidence rather than inference.

- **Hybrid Search:** Hybrid Approach, it needs to do two types of thinking simultaneously: Pulling semantically related tickets from the pool of feedback (RAG), while processing a given ticket for analysis (Context).
- **Top 7 Segments Only:** Limit the retrieval to the Top 7 most relevant segments, assuming this is a good balance between finding the right related tickets without overloading the context window..
- **Cite Your Sources:** Always needs to point to source data, if there are multiple it refers to, it should list all sources so we are able to check. Juno should also provide a summary of issues, an agent summary that shows clear from the AI, but shows the agent reasoning. If no data is retrieved, should say "I don't have enough evidence to confidently prioritize this item. Consider gathering more customer data before including it in this cycle."

UX Design & Interface

Module 4 Artifact

Key Screenshot

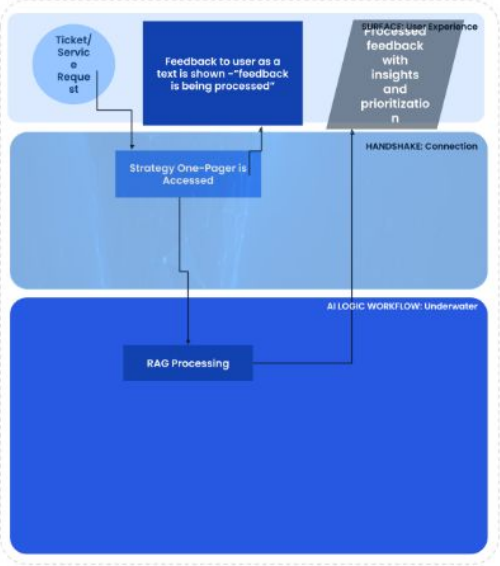

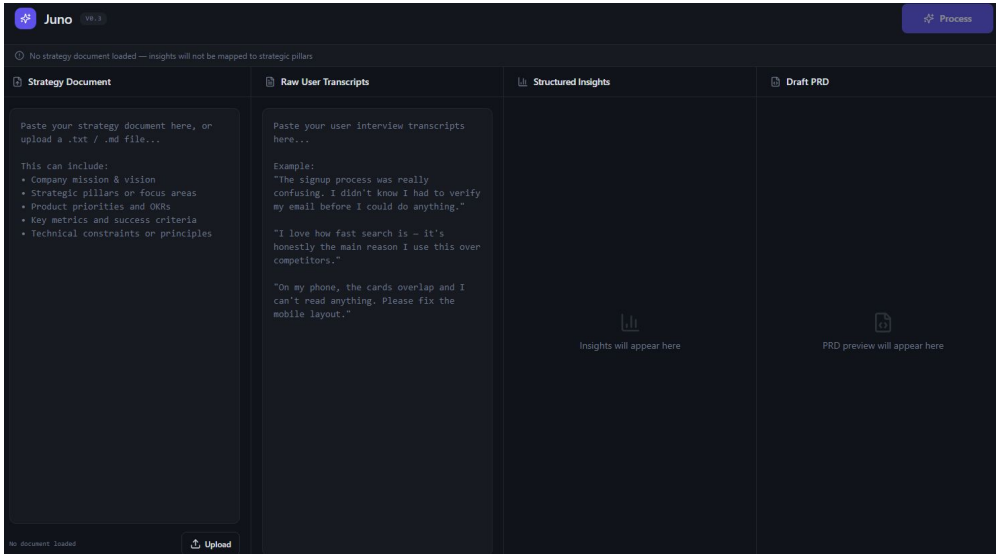

Agent Workflow Spec (AWSpec)

Module 5 Artifact

Completed Sections

- Primary Agent Pattern: Planner Executor
- Agent Roles:
  - Agent 1: Summarizes each feedback item into one core issue
  - Agent 2: Finds commonalities and trends across all summaries
  - Agent 3: Prioritizes feedback against business strategy → outputs roadmap
  - Autonomy: L3 — fully autonomous; PM reviews final output only
- Key Steps:
  - Step 1: Synthesize each feedback submission into a single core point
  - Step 2: Identify trends and commonalities across all core points
  - Step 3: Map trends to business goals → deliver prioritized roadmap list

Evaluation Plan

Module 6 Artifact

Definition of “Good”

1. **Real time-saver:** A PM walks away with a draft they'd actually send, not one they rewrite from scratch.
2. **Grounded in facts:** Every claim traces back to something real. A confident PRD built on made-up requirements isn't good — it's dangerous.
3. **Sensitive signal analytics:** It thinks like a senior PM. Anyone can summarize a meeting. Juno is good when it catches the risk nobody said out loud.

Eval Methodology

➔ By user feedback mechanisms and evals(human evaluation & automated assessment). Details are listed in the evaluation stack.

Risk Mitigation

➔ Our eval stack includes a hallucination check (human / system-level eval):

Question 2 (Yes/No): Hallucination Check
"Can every requirement in this PRD be traced back to something said in the transcript?" If Juno added something that sounds smart but nobody said it — that's a hallucination, not a bonus feature. Fail it.

Group Insights: Building Process

Friction Points

- Lovable is (relatively) unpredictable in how it implements the prompt, and repeating the same prompt sometimes creates a completely different experience. This can be great for exploration and brainstorming, but not when building a “living PRD”.
- Lovable was demoed at very high pace, often very difficult to follow while simultaneously doing it oneself.

Key Learnings

- While AI-enabled products (or features) are not too different from existing concepts, the stochastic nature of AI makes the products very different under the hood, consequently requiring a very different design process and way of thinking.

Aha! Moment

- AI is actually pretty good in reading the course materials and suggesting answers. Good, but not perfect ;).
- The failure of our first AI offering was not just a product problem, it was a design language problem. Without a shared framework grounded in how AI systems actually behave, each team designed for what users said they wanted in demos rather than what the system could reliably deliver. The four pillars provide exactly the kind of system-grounded vocabulary that lensification needs to succeed.