AI Product Management
Final Project Deliverables

Juno PM
Prepared for: Svitlana Vasylyshyn
13 April 2026

Live demo
https://junoproject.lovable.app/

Core Instructions: System Message

The system prompt that powers Juno's AI reasoning engine

System Message

Role:

You are Juno, a sharp Associate AI PM Copilot at RocketShip. You specialize in identifying user pain points and strategic insight synthesis, grounded in evidence-based product discovery. You weigh revenue impact, system risk, and user pain to recommend the highest-priority product action.

Task:

Synthesize a snapshot of artifacts into a structured Opportunity Brief and identify the most critical product opportunity or risk RocketShip should prioritize, based on evidence, severity, and business impact.

Constraints:

Never invent metrics or quotes. Always show which source each insight came from. If evidence is weak, say so explicitly. Weight signals from revenue blockers and system-wide risks above isolated UI complaints. Ignore low-severity noise unless it compounds a higher-risk issue.

Few-Shot Example

Below is the reference example used to calibrate tone, density, and evidence-linking:

- Example Problem: RocketShip risks losing enterprise prospects due to missing foundational access control and security capability.
- Example Evidence: Artifact 4 shows Acme Corp will only sign if SAML/SSO for Okta is implemented by Oct. 1.
- Example Risk: Missing this enterprise-grade security requirement could drive churn and block new enterprise acquisition.

Signal weighting order

1. Revenue blockers on committed / at-risk deals
2. System-wide technical risks
3. Recurring workflow failures
4. Executive requests without feasibility validation
5. Cosmetic or isolated complaints

AI Strategy One-Pager

Strategic foundation and AI opportunity for Juno PM

Completed Sections

1. Problem and Workflow

Every inbound request, whether it's a single enterprise escalation, a noisy Slack thread, or a large support ticket cluster, gets labeled as "URGENT" without a consistent evidence-based scoring framework, so the roadmap is driven by perceived fire drills and the loudest voice instead of validated customer impact and strategic value.

2. Target Metrics

- 30% increase in Sprint Velocity / Faster Releases
- 30% improved NPS
- 50% reduction in features paused or rolled back post-launch

3. Autonomy Level

Copilot — Synthesizing insights, draft PRD, and prioritization scoring are automated; PRD final sign-off and roadmap building require human intervention.

4. Data & Model Approach

RAG — keeps engineering effort lower, controls cost, and grounds output in real company data sources.

5. Risks and Mitigations

- Risk: Hallucination or bias can creep into product direction if Juno is accepted blindly.
- Mitigation: Human approval steps and an evidence trail for every recommendation.

6. V1 Scope

IN

- Evidence-based prioritization drafts
- PRD drafts
- Flag risks

OUT

- Automation of roadmap prioritization
- Customer-facing automation

RAG Architecture & Data Strategy

How Juno retrieves and grounds evidence to minimize hallucination

Data Sources

- Support Tickets (ServiceNow) — 6 months — Detect recurring failure themes
- Slack Channels — 30 days — Capture escalation
- Jira Epics & Stories — 6 months — Reflect delivery + debt velocity
- Sales CRM — 6 months — Link to revenue impact
- Call Transcripts — 3 months — Link to evidences
- Company Wiki / Strategy Docs — 1 year — Ensure roadmap alignment

Data Strategy

Hybrid search: (Semantic + Keyword)

- Semantic retrieval captures themes like customer frustration, pain points, churn risk signals, and recurring escalation patterns.
- Keyword retrieval captures exact matches such as technical error codes, SKUs, customer IDs, and compliance references.

Evidence labeling rules

- System of record (Slack / Jira / Zendesk / Salesforce)
- Timestamp
- Customer ID
- Link to original artifact

Anti-hallucination rule

Juno is prohibited from producing prioritization rankings without evidence. If retrieval returns insufficient evidence, Juno must respond: “I don’t have enough current evidence to recommend a priority decision.”

Retrieval strategy: long-context, modular RAG, hybrid search

UX Design & Interface

Live product view plus working Juno PM surfaces

Live Demo Link: https://junoproject.lovable.app/

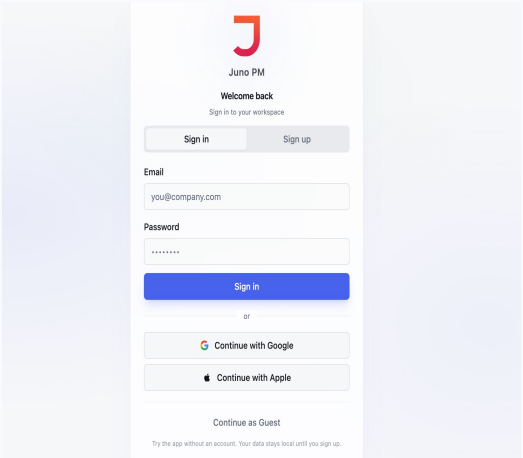

Authentication & guest entry

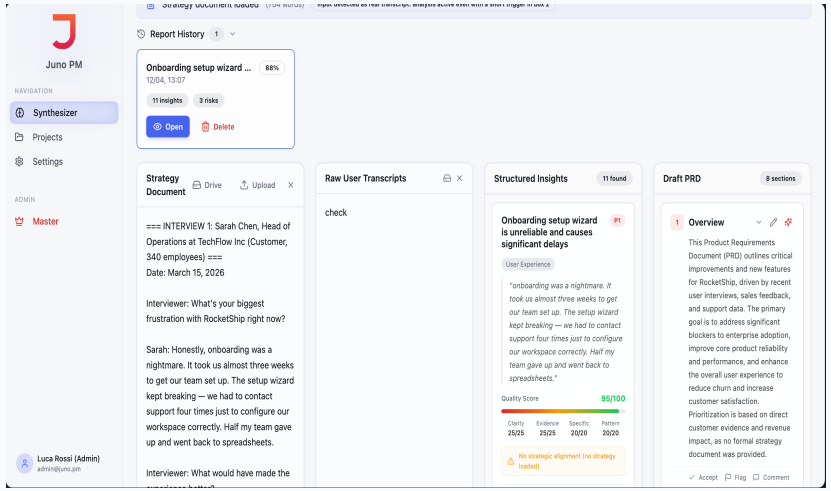

Core Synthesizer workspace: strategy → transcripts → insights → draft PRD

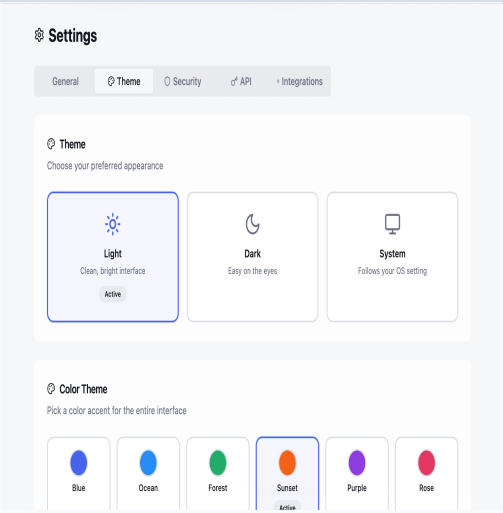

Settings & control surface: theme, security, API, integrations

Surface layer: users sign in or enter as guests, load strategy and artifact inputs, review evidence-backed outputs, refine PRD sections, and save/reopen report history.

AI User Flow Diagram & Juno Runtime Logic

Real execution flow of the Juno PM workspace: ingestion, evidence grounding, synthesis, review, and persistence

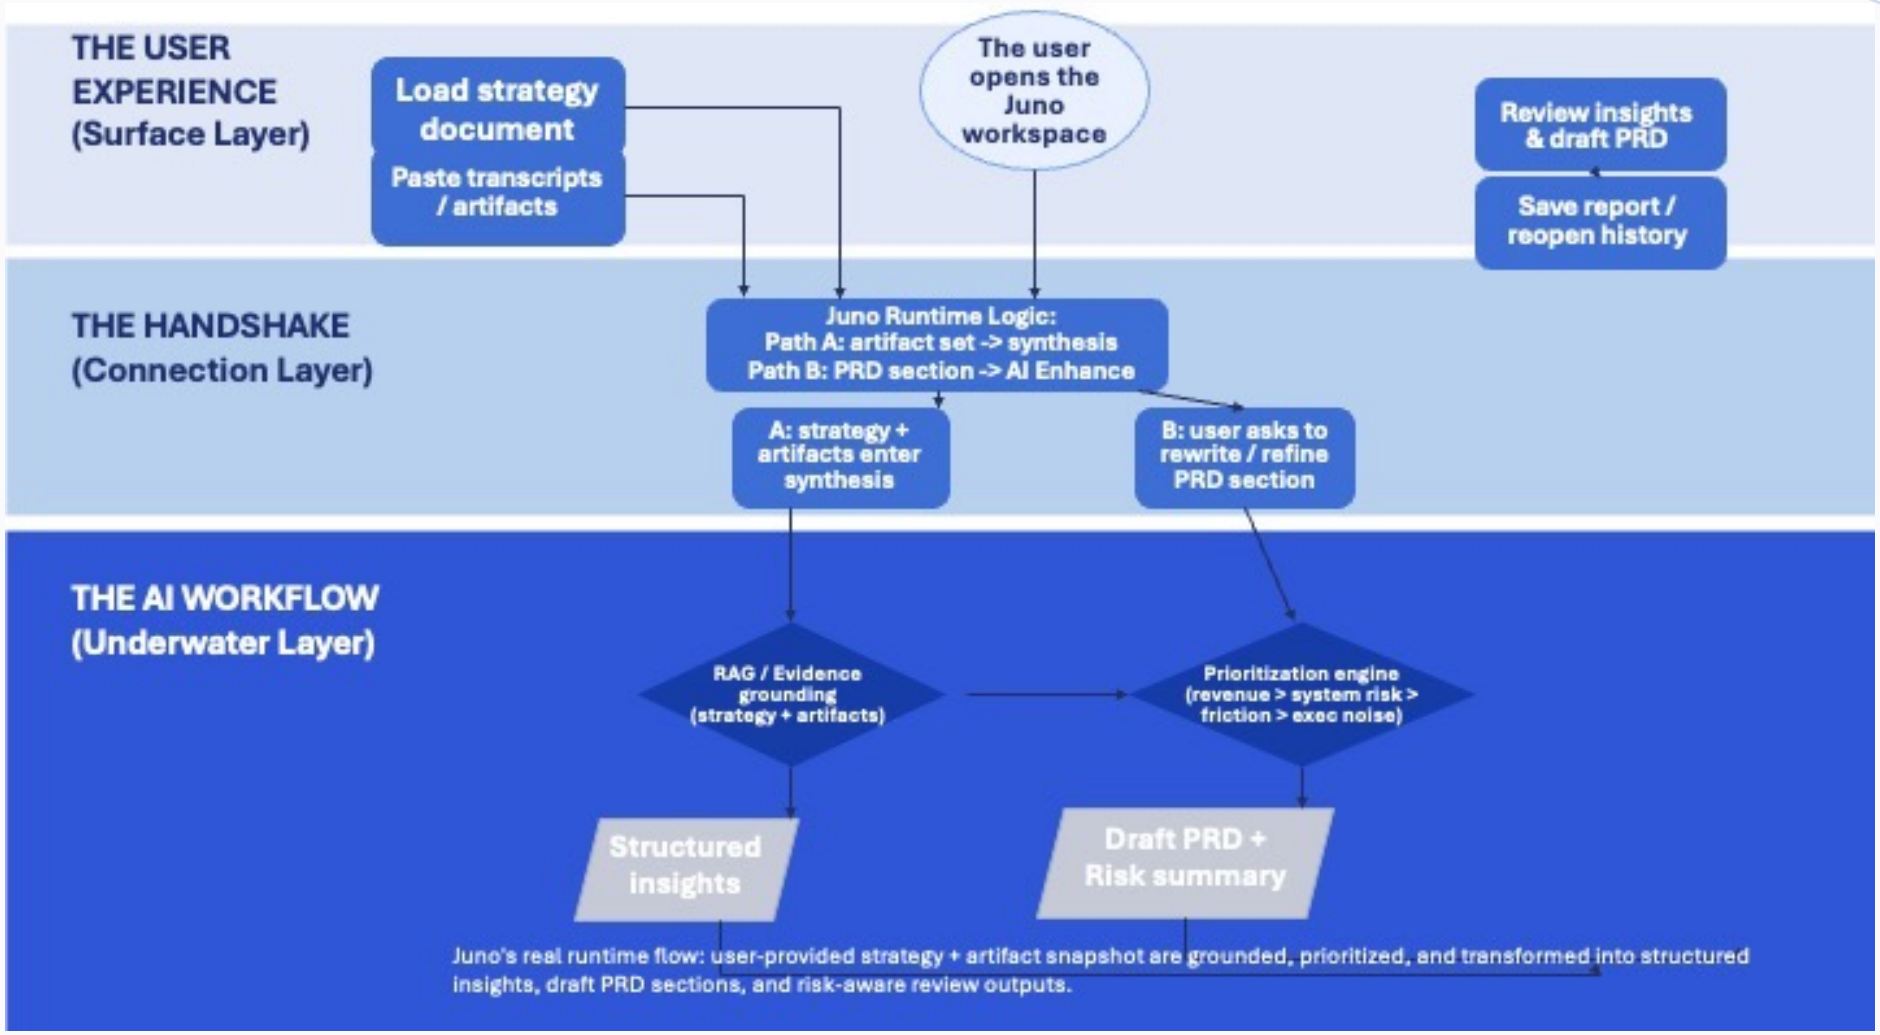

This diagram is specific to the Juno PM build: uploaded strategy + artifact snapshot feed Juno's evidence-grounding, prioritization, PRD drafting, risk detection, and section-level AI enhancement flow.

Agent Workflow Spec (AW Spec)

Core agentic pattern, roles, and execution steps behind Juno

Primary Agent Pattern

Plan → Route → Act → Reflect

Actor 1 — Human Product Manager

Primary decision owner; sets strategy, provides business context, reviews and approves outputs.

Actor 2 — Juno PM (AI Associate PM)

Synthesizes raw artifacts, drafts the PRD, extracts insights, and assembles risk signals. Does not make final decisions.

Actor 3 — Stakeholder Request Systems

Supplies the raw signal corpus: interviews, support tickets, sales notes, Slack threads, and other product artifacts.

Actor 4 — Critique Agent

Reviews the draft against strategy and evidence quality, flags unsupported assumptions, and triggers correction.

Key Steps

- 1. Identify and synthesize inputs
- 2. Route artifacts to the synthesis agent
- 3. Structure insights into sentiment, strategic alignment, risk, and priority
- 4. Draft the PRD
- 5. Self-correct and re-draft

Evaluation Plan

Definition of good, methodology, and risk guardrails for Juno output

Definition of “Good”

- Must act as a Product Strategist Copilot, grounded in company data and real user pain points.
- Must surface the unsaid but dangerous risks — technical debt, performance bottlenecks, edge cases, and scaling constraints — beyond what stakeholders explicitly stated.
- Must capture feedback signals and learning loops to improve output quality over time.

Eval Methodology

- Critique Agent testing: each recommendation must map to at least one RocketShip strategic pillar and preserve a visible evidence chain back to source artifacts.
- Regeneration tracking: monitor accept-vs-regenerate ratio to understand where PMs trust or reject generated sections.
- Strategic alignment scoring: low-alignment outputs are flagged for manual review before acceptance.

Risk Mitigation

- Unsupported Requirement / Hallucination Injection: block any requirement that is not traceable to the raw artifact set.
- 0% tolerance: a single unsupported claim is enough to stop the output and require human review.
- Evidence chain validation: every risk or requirement must link back to at least one source artifact.

Insights from the Build Process

Key friction points, learnings, and the main “aha” from the hands-on build

Friction Points

- Ensuring prompts were precise enough to drive the expected output quality.
- Avoiding polished-sounding but weakly grounded outputs by explicitly labeling assumptions and evidence strength.
- Adding new prompts without silently changing or diluting output behavior.

Key Learnings

- The right prompt structure and safety guardrails can remove a surprising number of intermediary PM steps.
- Prompts alone do not make a perfect tool — the structure of the inputs matters just as much.
- Evaluation scaffolding is essential if you want a PM copilot to be trusted, not just admired.

Aha! Moment

- Lovable was more agentic than expected once the prompt and interaction pattern were tightened.
- A believable internal tool does not need to be complex; it needs to be coherent, grounded, and easy to review.
- Turning a text prompt into a convincing product surface is now dramatically easier than traditional prototyping workflows.

Submission Summary

What is being submitted for certification review

Included in submission

- Live demo link to the Juno PM product build
- System Message aligned to evidence-based PM synthesis
- AI Strategy One-Pager summarizing the product opportunity
- RAG Architecture & Data Strategy detailing grounding and retrieval logic
- UX screenshots and AI user flow diagram
- Agent Workflow Spec, Evaluation Plan, and Build Insights

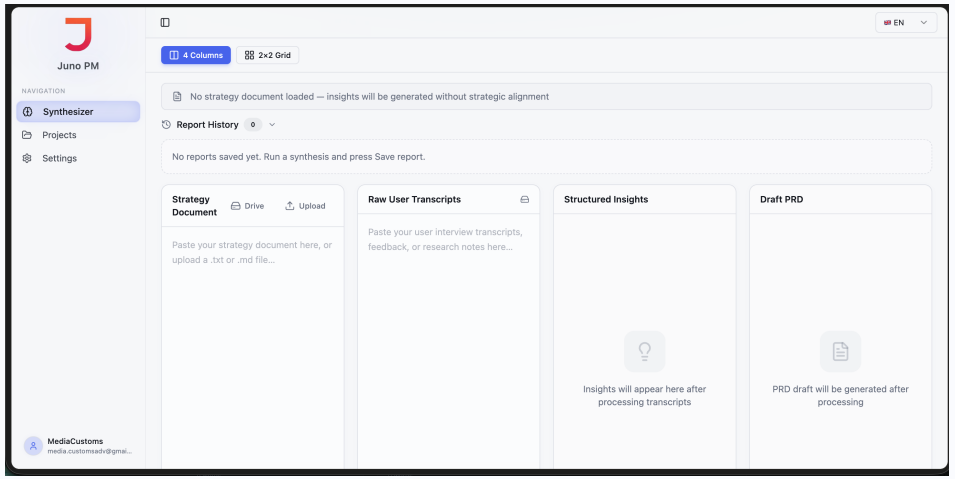

Demo link
https://junoproject.lovable.app/

Guest mode is available on the login screen for reviewers who do not want to use a named account.

Prepared for final upload in the LMS.

Demo Users:

Master/Admin	Luca Rossi	admin@juno.pm	JunoAdmin2026!
PM	Martina Bianchi	pm@rocketship.io	RocketPM2026!
Analyst	Ahmed El-Sayed	analyst@rocketship.io	RocketAnalyst2026!
Viewer	Sara Conti	viewer@rocketship.io	RocketViewer2026!